

Question:

Objective: The objective of this assignment is to investigate the performance of two different types of sentence-level representations as well as understand the type of information encoded in these representations as produced by the bert-base-uncased model.

There are usually two ways of aggregating sentence-level representation from the bert model.

- I. Using the CLS token
- II. Applying some transformation on the hidden state output of the model One such transformation is the application of max pooling.

In this assignment, you will be given a set of two sentences that will be taken as input from the system arguments with the sentences being separated “ , ”. You will have to use the pre-trained bert-base-uncased model to generate two different representations (as described above) for each sentence. You will then apply cosine similarity between the respective representations of each sentence and report the two values rounded off to the 2nd decimal.

The output of the assignment can be represented as,

$$\cos(\gamma_1(f(S_1)), \gamma_1(f(S_2))) \quad \cos(\gamma_2(f(S_1)), \gamma_2(f(S_2)))$$

Where $f(\cdot)$ is the bert-base-uncased model, S_1 and S_2 are the two sentences, γ_1 and γ_2 are the CLS and max-pooling representations of the output of the bert-base-uncased model and lastly, $\cos(\cdot, \cdot)$ is the cosine similarity score.

Bonus: The input cases have been provided such that the sentence similarities range from paraphrase to contradiction. Looking at the scores, try to understand which type of vector representation works best. Also, based on the last test case, do you think that the bert-baseuncase model is able to understand semantic information?

Sample Test Cases:

"input": "The brown fox jumped over the well , The brown fox did not jump over the well",

"output": "0.89 0.98"

"input": "The city of joy , better known as Kolkata",

"output": "0.51 0.76"